


# Advanced Section #1: SVMs, logistic regression and neural networks

CS 209B: Data Science

Javier Zazo

Pavlos Protopapas

[javier.zazo.ruiz@gmail.com](mailto:javier.zazo.ruiz@gmail.com)



# Lecture Outline

---

Classifying Linear Separable Data

Classifying Linear Non-Separable Data

Introduction to convex optimization

Dual problem of the SVC

Extension to Non-linear Boundaries


Relationship to Logistic Regression

## Classifying Linear Separable Data

---

# Decision Boundaries Revisited

- ▶ **Goal:** find *decision boundaries* to separate classes.
- ▶ Multiple decision boundaries:
  - Linear separable data → Hard margin classifier.
  - Linear non-separable data → SVCs.
  - Non-linear boundaries → SVMs.
- ▶ Logistic regression → Kernel Logistic Regression.
- ▶ Motivation for NN.



# Hyperplanes

- ▶ Hyperplane equation:

$$f(x) = \beta_0 + \beta^T x = 0.$$

- ▶ Normal vector:  $\beta$ .
- ▶ Halfspaces:  $f(x) \leq 0$ .

- ▶ Given  $N$  pair data points:


$$(x_1, y_1), \dots, (x_N, y_N)$$

- ▶ Binary classification:

$$y_i \in \{-1, 1\}$$

- ▶ Separable classes:


$$y_i(\beta_0 + \beta^T x_i) \geq 0$$



# Maximizing Margins

- ▶ We establish a geometric principle to classify data.
- ▶ Distance between  $x$  and hyperplane:

$$D(x) = \frac{\beta_0 + \beta^T x}{\|\beta\|}.$$



- ▶ Unsigned distances:  $|D(x_i)| = y_i D(x_i) \geq M$ .
- ▶ Maximum margin classifier:


$$\begin{aligned} & \max_{M, \beta_0, \beta} \quad M \\ & \text{s.t.} \quad \frac{1}{\|\beta\|} y_i (\beta_0 + \beta^T x_i) \geq M, \quad \forall i. \end{aligned} \tag{1}$$

# Maximum Margin Classifier

- We transform the previous problem into

$$\begin{aligned} \min_{\beta_0, \beta} \quad & \|\beta\| \\ \text{s.t.} \quad & y_i(\beta_0 + \beta^T x_i) \geq 1, \quad \forall i. \end{aligned} \tag{2}$$

- Both problems can be solved using specific solvers.
- They become unfeasible if the data is non-separable.




## Classifying Linear Non-Separable Data

---



# The Margin/Error Trade-Off

- ▶ With every decision boundary, bias-variance trade-off.
- ▶ The maximum margin classifier:
  - Low bias but **very high variance**.
- ▶ Generalization to linear non-separable boundaries.



# Support Vector Classifier (SVC)

---

- ▶ Classes overlap in the input space.
- ▶ Maximize the margin but relax the constraints.
- ▶ Use of slack variables  $\xi = (\xi_1, \dots, \xi_N)$ :

$$\begin{aligned} \min_{\beta_0, \beta, \xi_i \geq 0} \quad & \|\beta\| \\ \text{s.t.} \quad & y_i(\beta_0 + \beta^T x_i) \geq 1 - \xi_i, \quad \forall i \\ & \sum_{i=1}^N \xi_i \leq C, \end{aligned}$$

- ▶  $C$  limits the amount of violation of the constraints.
  - Margin violation: points inside the margin.
  - Misclassification: points on the wrong side of the boundary.

## Moving the constraint to the objective

---

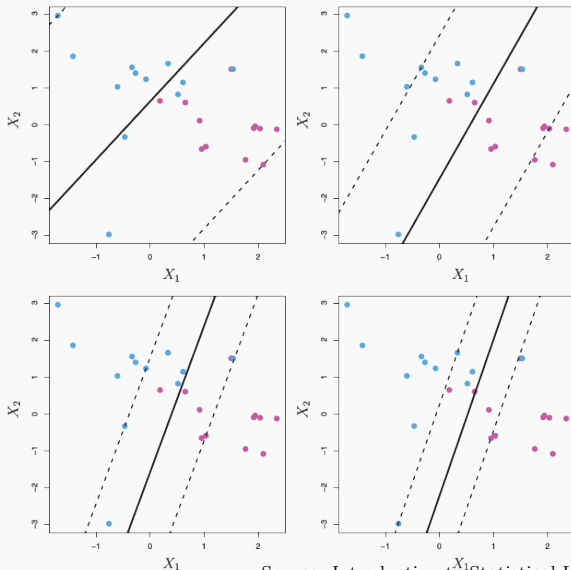
- Remove  $\sum_i \xi_i \leq C$  and put it in the objective:

$$\begin{aligned} \min_{\beta_0, \beta, \xi_i \geq 0} \quad & \frac{1}{2} \|\beta\|^2 + \lambda \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & y_i(\beta_0 + \beta^T x_i) \geq 1 - \xi_i, \quad \forall i. \end{aligned}$$

- Tuning SVCs:
  - small  $\lambda \rightarrow$  large margin
  - large  $\lambda \rightarrow$  narrow margins
  - $\lambda = \infty$  produces the hard margin solution

# Bias-variance trade-off examples

Lower to higher tuning parameter  $\lambda$ .



Source: Introduction to Statistical Learning

# Introduction to convex optimization

---

# Introduction to optimization


- Unconstrained optimization  $\rightarrow$  minimize in Euclidean space:

$$\min_{x \in \mathbb{R}^n} f(x). \quad (3)$$

- Constrained optimization  $\rightarrow$  minimization respect to  $X \subset \mathbb{R}^n$ :

$$\min_{x \in X} f(x). \quad (4)$$

- Convexity can be defined on sets and functions:



# Unconstrained optimization

---

- Necessary condition for point  $x^*$  be optimal:

$$\nabla_x f(x^*) = 0.$$

- **Convex:** condition is both necessary and sufficient.

- Example 1:  $f_1(x) = ax^2 + bx + c$  with  $a > 0$ .

$$- \frac{d}{dx} f(x) = 2ax + b = 0 \rightarrow \boxed{x^* = \frac{-b}{2a}}.$$

- Example 2:  $f_2(x) = x^T A x + 2b^T x + c$  and  $A \succ 0$ :

$$- \nabla_x f_2(x) = 2Ax + 2b = 0 \rightarrow \boxed{x^* = -A^{-1}b.}$$

- **Primal problem:** constrained problem in standard form:

$$\begin{array}{ll}\min_{x \in \mathbb{R}^n} & f(x) \\ \text{s.t.} & g_i(x) \leq 0 \quad i \in \{1, \dots, m\} \\ & h_j(x) = 0 \quad j \in \{1, \dots, p\}.\end{array}$$

- **Convex** if  $f(x)$ ,  $g_i(x)$  are convex, and  $h_j(x)$  are affine.
- **Slater's condition**  $\rightarrow$  establishes strong duality.
- **Duality theory** provides:
  - Optimality analysis.
  - Algorithmic tools.
  - Theoretical insights.



# Dual problem formulation

---

1. Construct the Lagrangian:

$$L(x, \lambda, \nu) = f(x) + \sum_i \lambda_i g_i(x) + \sum_j \nu_j h_j(x).$$

2. **Dual function:** the minimum of the Lagrangian over  $x$ :

$$q(\lambda, \nu) = \min_x L(x, \lambda, \nu).$$

3. **Dual problem:** maximization of the dual function over  $\lambda_i \geq 0$ :

$$\begin{aligned} \max_{\lambda \in \mathbb{R}^m, \nu \in \mathbb{R}^p} \quad & q(\lambda, \nu) \\ \text{s.t.} \quad & \lambda_i \geq 0 \quad \forall i. \end{aligned} \tag{5}$$

# Necessary conditions for optimality

---

Karush-Kuhn-Tucker (KKT) conditions:

- The minimization w.r.t.  $x$  of the Lagrangian is an unconstrained problem (step 2). Therefore,

$$\nabla_x L(x^*, \lambda, \nu) = 0 \quad (6)$$

is necessary for any candidate solution  $x^*$ .

- Feasibility is also required:

$$g_i(x^*) \leq 0 \quad \forall i \quad (7a)$$

$$h_j(x^*) = 0 \quad \forall j \quad (7b)$$

$$\lambda_i^* \geq 0 \quad \forall i \quad (7c)$$

$$\nu_j^* \in \mathbb{R} \quad \forall j, \quad (7d)$$

- These equations define a system to recover primal variables.

## Dual problem of the SVC

---

# Motivation for the dual problem

---

- ▶ **Primal problem:** linear classifier that should
  - maximize the distance between the points and the decision boundary (maximize margin).
  - misclassify as few points as possible.

$$\begin{aligned} \min_{\beta_0, \beta, \xi_i \geq 0} \quad & \frac{1}{2} \|\beta\|^2 + \lambda \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & y_i(\beta_0 + \beta^T x_i) \geq 1 - \xi_i, \quad \forall i. \end{aligned}$$

- ▶ **Dual problem:**
  - Allows for efficient computations (used in solvers).
  - Allows a natural derivation of kernel methods.
- ▶ **Consideration:** SVC is a convex problem.
  - Strong duality holds.

# Derivation of the dual problem

---

- Standard SVC:

$$\begin{aligned} \min_{\beta_0, \beta, \xi_i \geq 0} \quad & \frac{1}{2} \|\beta\|^2 + \lambda \sum_{i=1}^N \xi_i \\ \text{s.t.} \quad & y_i(\beta_0 + \beta^T x_i) \geq 1 - \xi_i, \quad \forall i \in \{1, \dots, N\}. \end{aligned}$$

- Lagrangian (it's a long expression):

$$L(\beta, \xi, \alpha, \mu) = f(\beta, \xi) + \sum_i \alpha_i g_i(\beta, \xi) + \mu_i \xi_i$$

- First order condition: gradient with respect to  $\beta$ ,  $\beta_0$  and  $\xi$

$$\beta = \sum_{i=1}^N \alpha_i y_i x_i, \quad 0 = \sum_{i=1}^N \alpha_i y_i, \quad \alpha_i = \lambda - \mu_i, \quad \forall i.$$

- With feasibility  $\rightarrow$  the system of equations can be solved exactly.

# Dual problem of SVCs

- ▶ Substituting on the primal problem we obtain the dual:

$$\begin{aligned} \max_{0 \leq \alpha_i \leq \lambda} \quad & \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j x_i^T x_j \\ \text{s.t.} \quad & \sum_{i=1}^N \alpha_i y_i = 0. \end{aligned}$$

- ▶ Sequential Minimal Optimization (SMO) solves the problem very efficiently.
- ▶ Only support vectors have  $\alpha_i \neq 0$ .
- ▶ At test time, we use

$$y_{\text{test}} \leftarrow \text{sign}[\beta_0 + \beta^T x_{\text{test}}].$$

## Extension to Non-linear Boundaries


# Transforming the Data

- ▶ Training a polynomial model is just training a linear model on data with transformed predictors:

$$\phi : \mathbb{R} \rightarrow \mathbb{R}^4$$

$$\phi(x) = (x^0, x^1, x^2, x^3)$$

- ▶  $\mathbb{R}$  is the *input space*;  $\mathbb{R}^4$  is the *feature space*.





# SVC with Non-Linear Decision Boundaries

## Generalization:

1. Apply transform  $\phi : \mathbb{R}^J \rightarrow \mathbb{R}^{J'}$  on training data

$$x_n \mapsto \phi(x_n)$$



where typically  $J'$  is much larger than  $J$ .

2. Train an SVC on the transformed data

$$\{(\phi(x_1), y_1), \dots, (\phi(x_N), y_N)\}$$

► XOR example:

$$\phi(x) = (x_1, x_2, x_1x_2)$$



# Training in the feature space

---

- We can train the SVC in the feature space:

$$\begin{aligned} \max_{0 \leq \alpha_i \leq \lambda} \quad & \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j \phi(x_i)^T \phi(x_j) \\ \text{s.t.} \quad & \sum_{i=1}^N \alpha_i y_i = 0. \end{aligned}$$

- Designing  $\phi(x)$  can be hard  $\rightarrow$  (maybe not with NN).
- Computing  $\phi$  explicitly can be costly.
- We are only interested in computing  $\phi(x_i)^T \phi(x_j)$ .

## Definition of Kernel

The *inner product* between two vectors is a measure of the similarity of the two vectors.

### Definition

Given a transformation  $\phi : \mathbb{R}^J \rightarrow \mathbb{R}^{J'}$ , from input space  $\mathbb{R}^J$  to feature space  $\mathbb{R}^{J'}$ , the function  $K : \mathbb{R}^J \times \mathbb{R}^J \rightarrow \mathbb{R}$  defined by

$$K(x_i, x_j) = \phi(x_i)^T \phi(x_j), \quad x_i, x_j \in \mathbb{R}^J$$

is called the *kernel function* of  $\phi$ .

Generally, *kernel function* may refer to any function  $K : \mathbb{R}^J \times \mathbb{R}^J \rightarrow \mathbb{R}$  that measure the similarity of vectors in  $\mathbb{R}^J$ , without explicitly defining a transform  $\phi$ .

- For a choice of kernel  $K(x_i, x_j) = \phi(x_i)^T \phi(x_j)$ ,

$$\begin{aligned} \max_{0 \leq \alpha_i \leq \lambda} \quad & \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N y_i y_j \alpha_i \alpha_j K(x_i, x_j) \\ \text{s.t.} \quad & \sum_{i=1}^N \alpha_i y_i = 0. \end{aligned}$$

without computing the mappings  $\phi(x_i), \phi(x_j)$ .

- This way of training a SVC in feature space without explicitly using the mapping  $\phi$  is called *the kernel trick*.

Common kernel functions include:

- ▶ **Polynomial Kernel**

$$K(x_1, x_2) = (x_1^\top x_2 + 1)^d$$

where  $d$  is a hyperparameter

- ▶ **Radial Basis Function Kernel**

$$K(x_1, x_2) = \exp \left\{ -\frac{\|x_1 - x_2\|^2}{2\sigma^2} \right\}$$

where  $\sigma$  is a hyperparameter

- ▶ **Sigmoid Kernel**

$$K(x_1, x_2) = \tanh(\kappa x_1^\top x_2 + \theta)$$

where  $\kappa$  and  $\theta$  are hyperparameters.

## Relationship to Logistic Regression

---

- General structure for *classification* and *regression*:

$$\min_{\beta} J(X, y, \beta) + \lambda P(\beta),$$

- $J(X, y, \beta)$  is a general loss function.
- $P(\beta)$  is a general regularizer.

- SVCs take a similar form as well:

$$\min_{\beta, \beta_0} \sum_{i=1}^N \max[0, 1 - y_i(\beta_0 + \beta^T x_i)] + \lambda \|\beta\|^2,$$

$$\begin{aligned} y_i(\beta_0 + \beta^T x) &\geq 1 - \xi_i \rightarrow \xi_i = 1 - y_i(\beta_0 + \beta^T x) \geq 0 \\ \rightarrow \xi_i &= \max[0, 1 - y_i(\beta_0 + \beta^T x)]. \end{aligned}$$

# SVMs via Loss + Penalty

We can extend SVCs to incorporate feature maps  $\phi$ :

$$\min_{\beta, \beta_0} \sum_{i=1}^N \max[0, 1 - y_i(\beta_0 + \beta^T \phi(x_i))] + \lambda \|\beta\|^2.$$

Notice that from KKT condition,  $\beta = \sum_{j=1}^N y_j \alpha_j \phi(x_j)$ :

$$\min_{\beta, \beta_0} \sum_{i=1}^N \max[0, 1 - y_i(\beta_0 + \sum_{j=1}^N y_j \alpha_j \phi(x_j)^T \phi(x_i))] + \lambda \|\beta\|^2.$$

and introducing kernel  $K(x_i, x_j) = \phi(x_i)^T \phi(x_j)$

$$\min_{\beta, \beta_0} \sum_{i=1}^N \max[0, 1 - y_i(\beta_0 + \sum_{j=1}^N y_j \alpha_j K(x_i, x_j))] + \lambda \|\mathbf{K}^{1/2} \alpha\|^2.$$



# SVMs via Loss + Penalty and Kernels

- General cost function for SVMs:

$$\max[0, 1 - yf(x)] + \lambda \|\mathbf{K}^{1/2}\alpha\|^2$$

with

$$f(x) = \beta_0 + \sum_{j=1}^N \alpha_j K(x, x_j)$$

- Our training variables are now  $\beta_0$  and  $\alpha_j$ .
- We perform classification as usual:

$$y_{\text{test}} \leftarrow \text{sign}[f(x)].$$

# Logistic regression

---

- ▶ Given training data  $x_i \in \mathbb{R}^p$ ,  $y_i \in \{1, 0\}$ .
- ▶ Probabilities  $\rightarrow$  based on the logistic function:

$$p = P(y = 1|x) = \frac{1}{1 + e^{-(\beta_0 + \beta^T x)}},$$

- ▶ Likelihood  $\rightarrow$  Bernoulli distribution:

$$J_l(X, y, \beta) = \prod_{i=1}^N P(y = y_i|x) = \prod_{i=1}^N p^{y_i} (1 - p)^{1-y_i}$$

- ▶ Objective  $\rightarrow$  maximum log-likelihood:

$$\max_{\beta_0, \beta} \log(J_l(X, y, \beta)) \iff$$

$$\min_{\beta_0, \beta} - \sum_{i=1}^N [y_i \log(1 + e^{-(\beta_0 + \beta^T x_i)}) + (1 - y_i) \log(1 + e^{(\beta_0 + \beta^T x_i)})]$$

# Kernel Logistic Regression (KLR)

- ▶ Logistic regression  $\rightarrow$  loss function + penalization.
- ▶ We incorporate  $f(x)$  based on **kernels**:

$$\min_{\beta_0, \alpha_i} - \sum_{i=1}^N [y_i \log(1 + e^{-f(x_i)}) + (1 - y_i) \log(1 + e^{f(x_i)})] + P(\alpha),$$

$$f(x) = \beta_0 + \sum_{j=1}^N \alpha_j K(x, x_j).$$



# Comparison between SVMs and KLR

---

- ▶ Classification performance is very similar.
- ▶ KLR provides estimates of class probabilities.
- ▶ KLR generalizes naturally to M-class classification.
- ▶ KLR converges to the maximum margin classifier.
- ▶ KLR is computationally more expensive,  $O(N^3)$  vs.  $O(N^2m)$ .
- ▶ In SVMs many  $\alpha_i$  are zero  $\rightarrow$  data compression.
- ▶ In KLR all  $\alpha_i$  are typically non-zero.

# Logistic regression as Neural Network

- ▶ Logistic regression constitutes a NN of a single neuron.
- ▶ Sigmoid function:  $\sigma(z) = 1/(1 + \exp(-z))$ .

